



Claude Certifications

The complete field guide to the four Anthropic Claude certifications: what each exam measures, what it costs, how to prepare for it, and how to sit it.

Associate Foundations

Developer Foundations

Architect Foundations

Architect Professional

320

practice questions

110

flashcards

22

courses covered

33

pages



Compiled by Amey Thakur

Written while preparing for these exams, after working through every course in the official curriculum.

github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS · amey-thakur.github.io/CLAUDE-CERTIFICATIONS

Facts drawn from the official Anthropic exam guides and program pages. A community study resource, not affiliated with or endorsed by Anthropic. Free to share.

Why this companion exists

I sat down to prepare for these certifications and found the material scattered: the blueprint in one PDF, the policies in another, the courses on a platform behind a sign-in, the registration mechanics somewhere else again, and the practical details, the ones that actually cost people a sitting, written down nowhere at all.

So I collected it. Every official exam guide, every policy, every course, the blueprints with their real weights, and the things I only learned by going through it: which domains carry the marks, which options an exam question is quietly steering you away from, and how much of a failed attempt is logistics rather than knowledge.

None of this required anything you do not already have. The official material is free, the blueprints tell you exactly what is tested, and the rest is steady work.

What I could not do is give you the questions, and I would not if I could. Every candidate signs a non-disclosure agreement, and a credential is only worth holding if the people holding it earned it. So everything here is built from published material: the official blueprints, the published scenarios, and practice written from scratch against them.

Read the roadmap, pick the exam that matches the work you already do, and work its page. Keep the cheat sheet for the day before. If a fact here disagrees with the official exam guide, believe the guide, and please tell me so I can fix it.

I put this in one place so your time goes into learning rather than looking. If it helps you get certified, it did its job. Good luck.

– Amey Thakur



What is in this companion

Five parts, in the order a candidate needs them: choose the exam, learn what it measures, prepare for it, sit it, and know where to go afterwards.

PART 1	Choose your exam	Four certifications across three roles, with no prerequisites in either direction. Pick the one that matches the work you already do.	5
PART 2	Know your exam	A page of facts and a cheat sheet for each certification: the blueprint with real domain weights, and the rules that decide questions.	8
PART 3	Prepare	What to study, in what order, and how to practice without touching material you are not allowed to touch.	18
PART 4	Sit it	Booking, proctoring, and the day itself, then what the result means and what happens if it is not the one you wanted.	25
PART 5	Keep going	You now have everything the exam asks of you. What follows is where to find the parts that keep moving.	29

How to use it. If you already know which exam you are taking, go straight to its pages in Part 2 and keep the cheat sheet for the day before. Everything here is derived from the official exam guides, which remain the authoritative source and are mirrored in the repository.



What I would tell myself before the first exam

Ten things that are all published somewhere, and that almost nobody reads until it is too late to act on them. If you take one page from this companion, take this one.

The blueprint is the scope, and it is enforced. Every exam guide says it plainly: anything outside the blueprint is not tested. People lose weeks reading around the subject. Print the blueprint, mark each objective red, amber, or green, and study red first.

Study by weight, not by interest. Applications and integration is 33.1% of the Developer paper. Claude Code is 3.1%. Two evenings on the wrong one is a real cost. The weights are published; use them as your timetable.

A failed attempt costs weeks, not just money. Fourteen days before a second attempt, thirty before a third, ninety before a fourth, and four attempts maximum in a rolling twelve months. That waiting period, not the fee, is what hurts.

Book earlier than you feel ready. Rescheduling is free until twenty-four hours before. An early booking costs nothing and supplies the deadline that makes the studying happen. Inside twenty-four hours the fee is gone.

Most failed sittings are logistics. Run the OnVUE system test on the machine and the network you will actually use, not a different one. A corporate network that blocks a single domain will end the exam after you have paid for it.

Fix your partner email before you plan anything. Domain record changes take seven to ten days. Discovering this the week you meant to sit is the most avoidable delay in the whole process.

For Architect Foundations, rehearse all six scenarios. The exam presents four, drawn from a published set of six. They are not a surprise. Working through all six beforehand converts the hardest paper into a familiar one.

The credential expires in twelve months. Renewal is free, open book, and not proctored, so it is easy. Letting it lapse is not: you sit the full exam again at full price. Put the date in a calendar the day you pass.

Read the score report properly if you fail. It gives percent-correct by domain. Weight each by its share of the paper and you will usually find the gap sits in one or two places, not everywhere.

Closed book means closed book. No notes, no documentation, no translation tools, no Claude. Everything in this companion is for before the exam, not during it.



PART 1

Choose your exam

Four certifications across three roles, with no prerequisites in either direction. Pick the one that matches the work you already do.

IN THIS PART

Which certification is yours?

6

The roadmap

7



Which certification is yours?

Pick the exam that matches the work you already do.

 **START HERE**

Which certification is yours?

Pick the exam that matches the work you already do. There are no prerequisites, and each credential lasts 12 months.

You advise customers and run engagements
Associate · Foundations
CCAO-F · 60 items · \$99
Does not count toward partner tier

- Turning business problems into Claude workflows
- Judging output quality and knowing when to escalate
- Data sensitivity and responsible use

You build with the API, Claude Code, or MCP
Developer · Foundations
CCDV-F · 53 items · \$125
Counts toward partner tier

- API mechanics: messages, streaming, batches, caching
- Agents, tools, MCP servers, and structured output
- Security, cost control, and evaluation

You design Claude solutions end to end
Architect · Foundations
CCAR-F · 60 items · \$125
Counts toward partner tier

- Agentic architecture and orchestration
- Claude Code configuration for teams and CI
- Four of six published scenarios appear on your paper

You govern Claude systems at enterprise scale
Architect · Professional
CCAR-P · 63 items · \$175
No prerequisite. Foundations does not upgrade into it

- Integration, retrieval, and observability at scale
- Evaluation frameworks and governance
- Communicating architecture to stakeholders

Every exam: 120 minutes, closed book, proctored by Pearson VUE, passing at 720 of 1000, valid 12 months with a free renewal assessment.

 **Amey Thakur**
github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS

ANTHROPIC
Facts drawn from the official Anthropic exam guides



The roadmap

All four certifications: what each measures, what it costs, and what prepares you for it.



Claude Certification Roadmap

Four certifications, three roles. What each exam covers, what it costs, and what prepares you for it.

ASSOCIATE

Foundations

CCAO-F · 60 items · \$99
120 min · pass 720/1000

Consultants, sellers, and delivery leads

DOMAINS BY WEIGHT

Output evaluation	21%
Workflow integration	16%
Governance and risk	15%
Prompting	14%
Product and model choice	12%
Configuration	12%
Troubleshooting	10%

PREPARE WITH

Claude 101
Claude Platform 101
AI Fluency: Framework
AI Capabilities and Limits
Introduction to Cowork

Does not count toward partner tier

DEVELOPER

Foundations

CCDV-F · 53 items · \$125
120 min · pass 720/1000

Engineers building on the Claude platform

DOMAINS BY WEIGHT

Applications and integration	33.1%
Model selection and optimization	16.8%
Agents and workflows	14.7%
Prompt and context engineering	11.0%
Tools and MCPs	10.6%
Security and safety	8.1%
Claude Code	3.1%
Eval, testing, and debugging	2.6%

PREPARE WITH

Building with the Claude API
Introduction to MCP
Claude Code in Action
Introduction to agent skills
Introduction to subagents

Counts toward partner tier

ARCHITECT

Foundations

CCAR-F · 60 items · \$125
120 min · 4 of 6 scenarios

Architects building production systems

DOMAINS BY WEIGHT

Agentic architecture	27%
Claude Code workflows	20%
Prompting, structured output	20%
Tool design and MCP	18%
Context and reliability	15%

PREPARE WITH

Claude Code in Action
Introduction to MCP
Building with the Claude API
Agent skills and subagents
The six published scenarios

Counts toward partner tier

ARCHITECT

Professional

CCAR-P · 63 items · \$175
120 min · pass 720/1000

Architects governing systems at scale

DOMAINS BY WEIGHT

Integration	19%
Solution design	17%
Evaluation and testing	16%
Governance and risk	14%
Stakeholders, lifecycle	14%
Models and prompting	13%
Developer enablement	7%

PREPARE WITH

MCP: Advanced Topics
Claude with Amazon Bedrock
Claude on Google Cloud
An end-to-end system you built
Anthropic engineering articles

No prerequisite. Foundations does not upgrade

Every exam: 120 minutes · closed book · Pearson VUE, online or test center · pass at 720 of 1000 · valid 12 months, free renewal · badge via Credly · partner tiers get 50% off



Amey Thakur

github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS

ANTHROPIC

Facts drawn from the official Anthropic exam guides



PART 2

Know your exam

A page of facts and a cheat sheet for each certification: the blueprint with real domain weights, and the rules that decide questions.

IN THIS PART

Claude Certified Associate, Foundations	9
Cheat sheet: Associate Foundations	10
Claude Certified Developer, Foundations	11
Cheat sheet: Developer Foundations	12
Claude Certified Architect, Foundations	13
Cheat sheet: Architect Foundations	14
The six published scenarios	15
Claude Certified Architect, Professional	16
Cheat sheet: Architect Professional	17



ASSOCIATE · FOUNDATIONS

Claude Certified Associate

CCAO-F

exam code

60

items

120

minutes

720

to pass

\$99

list fee

12

months valid

WHERE THE MARKS ARE

Output evaluation	21%
Workflow integration	16%
Governance and risk	15%
Prompting	14%
Product and model choice	12%
Configuration	12%
Troubleshooting	10%

WHO IT IS FOR

Consultants, sellers, and delivery leads. Does not count toward partner tier.

IF YOU REMEMBER NOTHING ELSE

1. Verify anything specific and consequential. Confidence is not evidence.
2. Omission fails too. Read the source, not only the output.
3. Anonymize, then analyze. Instructions are not a control.
4. Match the model to the task in both directions.
5. Structure beats intensifiers: role, sections, constraints, example.
6. Projects hold what repeats; prompts hold what changes.
7. Project knowledge is maintained by you, not self-updating.
8. Diagnose what changed before rewriting anything.
9. Escalate system integrations to Developer and Architect scope.

PREPARE WITH

Claude 101 · Claude Platform 101 · AI Fluency: Framework · AI Capabilities and Limits · Introduction to Cowork



Cheat sheet: Associate Foundations

The page worth keeping for the hour before the exam.

 **ASSOCIATE · FOUNDATIONS · CHEAT SHEET**
Claude Certified Associate

CCAO-F exam code	60 items	120 minutes	720 to pass	\$99 list fee	12 months valid
---	-----------------------------------	--------------------------------------	--------------------------------------	--	--

WHERE THE MARKS ARE

Output evaluation	21%
Workflow integration	16%
Governance and risk	15%
Prompting	14%
Product and model choice	12%
Configuration	12%
Troubleshooting	10%

IF YOU REMEMBER NOTHING ELSE

- 1 Verify anything specific and consequential. Confidence is not evidence.
- 2 Omission fails too. Read the source, not only the output.
- 3 Anonymize, then analyze. Instructions are not a control.
- 4 Match the model to the task in both directions.
- 5 Structure beats intensifiers: role, sections, constraints, example.
- 6 Projects hold what repeats; prompts hold what changes.
- 7 Project knowledge is maintained by you, not self-updating.
- 8 Diagnose what changed before rewriting anything.
- 9 Escalate system integrations to Developer and Architect scope.

ON THE DAY

Answer everything, unanswered scores zero · flag and return rather than stalling · read how many responses each item wants · find the stated constraint, then eliminate

Full sheet, study guide, notes, practice questions, and a timed mock exam: amey-thakur.github.io/CLAUDE-CERTIFICATIONS

 **Amey Thakur**
github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS

ANTHROPIC
Facts drawn from the official Anthropic exam guide

The full sheet, with the reasoning behind every rule and a timed mock exam for this certification, is at amey-thakur.github.io/CLAUDE-CERTIFICATIONS



Claude Certified Developer

CCDV-F

exam code

53

items

120

minutes

720

to pass

\$125

list fee

12

months valid

WHERE THE MARKS ARE

Applications and integration	33.1%
Model selection and optimization	16.8%
Agents and workflows	14.7%
Prompt and context engineering	11.0%
Tools and MCPs	10.6%
Security and safety	8.1%
Claude Code	3.1%
Eval, testing, and debugging	2.6%

WHO IT IS FOR

Engineers building on the Claude platform. Counts toward partner tier.

IF YOU REMEMBER NOTHING ELSE

1. stop_reason drives the loop: tool_use executes, end_turn finishes.
2. Batches for latency-tolerant volume: half cost, 24-hour window.
3. Stable prefix first, then cache it. Cuts latency and cost together.
4. Enforce a schema, validate, retry with the validation error.
5. Known path is a workflow. Unknown path earns an agent.
6. Subagents exist to isolate context, not to add horsepower.
7. Pin model versions; upgrades become evaluated changes.
8. Injection is defeated structurally, never by asking politely.
9. A tool the agent does not hold cannot be misused.

PREPARE WITH

Building with the Claude API · Introduction to MCP · Claude Code in Action · Introduction to agent skills · Introduction to subagents



Cheat sheet: Developer Foundations

The page worth keeping for the hour before the exam.

 **DEVELOPER · FOUNDATIONS · CHEAT SHEET**
Claude Certified Developer

CCDV-F exam code	53 items	120 minutes	720 to pass	\$125 list fee	12 months valid
---	-----------------------------------	--------------------------------------	--------------------------------------	---	--

WHERE THE MARKS ARE

Applications and integration	33.1%
Model selection and optimization	16.8%
Agents and workflows	14.7%
Prompt and context engineering	11.0%
Tools and MCPs	10.6%
Security and safety	8.1%
Claude Code	3.1%
Eval, testing, and debugging	2.6%

IF YOU REMEMBER NOTHING ELSE

- 1 stop_reason drives the loop: tool_use executes, end_turn finishes.
- 2 Batches for latency-tolerant volume: half cost, 24-hour window.
- 3 Stable prefix first, then cache it. Cuts latency and cost together.
- 4 Enforce a schema, validate, retry with the validation error.
- 5 Known path is a workflow. Unknown path earns an agent.
- 6 Subagents exist to isolate context, not to add horsepower.
- 7 Pin model versions; upgrades become evaluated changes.
- 8 Injection is defeated structurally, never by asking politely.
- 9 A tool the agent does not hold cannot be misused.

ON THE DAY

Answer everything, unanswered scores zero · flag and return rather than stalling · read how many responses each item wants · find the stated constraint, then eliminate

Full sheet, study guide, notes, practice questions, and a timed mock exam: amey-thakur.github.io/CLAUDE-CERTIFICATIONS

 **Amey Thakur**
github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS

ANTHROPIC
Facts drawn from the official Anthropic exam guide

The full sheet, with the reasoning behind every rule and a timed mock exam for this certification, is at amey-thakur.github.io/CLAUDE-CERTIFICATIONS



Claude Certified Architect

CCAR-F

exam code

60

items

120

minutes

720

to pass

\$125

list fee

12

months valid

WHERE THE MARKS ARE

Agentic architecture 27%

Claude Code workflows 20%

Prompting, structured output 20%

Tool design and MCP 18%

Context and reliability 15%

WHO IT IS FOR

Architects building production systems. Counts toward partner tier.

IF YOU REMEMBER NOTHING ELSE

1. Agent loops need engineered termination, not a token budget.
2. Escalate on policy gaps and lack of progress.
3. Structured errors: a category and a retryable flag.
4. Tool descriptions are the selection mechanism. Differentiate them.
5. Hard limits belong in the tool layer, never in a prompt.
6. CLAUDE.md hierarchy, with path-scoped rules for conditional context.
7. CI means headless: print mode, JSON output, a schema.
8. Subagents return findings; raw documents stay at the edge.
9. Provenance must survive every handoff.

PREPARE WITH

Claude Code in Action · Introduction to MCP · Building with the Claude API · Agent skills and subagents · The six published scenarios



Cheat sheet: Architect Foundations

The page worth keeping for the hour before the exam.

 ARCHITECT · FOUNDATIONS · CHEAT SHEET
Claude Certified Architect

CCAR-F exam code	60 items	120 minutes	720 to pass	\$125 list fee	12 months valid
---	-----------------------------------	--------------------------------------	--------------------------------------	---	--

WHERE THE MARKS ARE

Agentic architecture	27%
Claude Code workflows	20%
Prompting, structured output	20%
Tool design and MCP	18%
Context and reliability	15%

IF YOU REMEMBER NOTHING ELSE

- 1 Agent loops need engineered termination, not a token budget.
- 2 Escalate on policy gaps and lack of progress.
- 3 Structured errors: a category and a retryable flag.
- 4 Tool descriptions are the selection mechanism. Differentiate them.
- 5 Hard limits belong in the tool layer, never in a prompt.
- 6 CLAUDE.md hierarchy, with path-scoped rules for conditional context.
- 7 CI means headless: print mode, JSON output, a schema.
- 8 Subagents return findings; raw documents stay at the edge.
- 9 Provenance must survive every handoff.

ON THE DAY

Answer everything, unanswered scores zero · flag and return rather than stalling · read how many responses each item wants · find the stated constraint, then eliminate

Full sheet, study guide, notes, practice questions, and a timed mock exam: amey-thakur.github.io/CLAUDE-CERTIFICATIONS

 **Amey Thakur**
github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS

ANTHROPIC
Facts drawn from the official Anthropic exam guide

The full sheet, with the reasoning behind every rule and a timed mock exam for this certification, is at amey-thakur.github.io/CLAUDE-CERTIFICATIONS



The six published scenarios

Four of these six frame every question on this paper. Rehearse them.

 ARCHITECT · FOUNDATIONS

The six published exam scenarios

Four of these six frame every question on your paper. Knowing them in advance is the largest preparation advantage on any Claude exam.

1

Customer support resolution agent

Agent SDK with MCP tools: `get_customer`, `lookup_order`, `process_refund`, `escalate_to_human`
Loop termination · escalation triggers · tool scoping · structured errors

2

Code generation with Claude Code

A team using Claude Code for generation, refactoring, debugging, and documentation
CLAUDE.md hierarchy · path-scoped rules · plan mode · custom commands

3

Multi-agent research system

A coordinator delegating to search, analysis, synthesis, and report subagents
Context handoffs · provenance · error propagation · citation integrity

4

Developer productivity tooling

Agents exploring unfamiliar codebases with built-in tools and MCP servers
Read, Write, Bash, Grep, Glob selection · MCP integration · large-context strategy

5

Claude Code in CI/CD

Automated code review, test generation, and pull request feedback
Headless invocation · JSON output · explicit review criteria · false positives

6

Structured data extraction

Extraction from unstructured documents, validated against JSON schemas
Schema design · validation retry loops · batch processing · confidence scoring

Rehearse each scenario until its likely questions are predictable. Every scenario names its primary domains in the official exam guide.

 **Amey Thakur**
github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS

ANTHROPIC
Scenarios published in the official Anthropic exam guide

Practice questions framed inside these scenarios are in the practice engine at amey-thakur.github.io/CLAUDE-CERTIFICATIONS



Claude Certified Architect

CCAR-P

exam code

63

items

120

minutes

720

to pass

\$175

list fee

12

months valid

WHERE THE MARKS ARE

Integration	19%
Solution design	17%
Evaluation and testing	16%
Governance and risk	14%
Stakeholders, lifecycle	14%
Models and prompting	13%
Developer enablement	7%

WHO IT IS FOR

Architects governing systems at scale. No prerequisite.
Foundations does not upgrade.

IF YOU REMEMBER NOTHING ELSE

1. Climb the pattern ladder only when the rung below cannot hold.
2. Least privilege means removal, not monitoring.
3. Retrieval degradation after a data change is an index problem.
4. Subjective quality is measured with rubrics and labeled sets.
5. Prompt changes are production changes: regression, then A/B.
6. Compliance is enforced before data crosses the boundary.
7. Route human review by confidence and consequence.
8. Present segmented, honest numbers, never a flattering average.
9. Observability captures traces, not just outputs.

PREPARE WITH

MCP: Advanced Topics · Claude with Amazon Bedrock · Claude on Google Cloud · An end-to-end system you built · Anthropic engineering articles



Cheat sheet: Architect Professional

The page worth keeping for the hour before the exam.

 ARCHITECT · PROFESSIONAL · CHEAT SHEET
Claude Certified Architect

CCAR-P exam code	63 items	120 minutes	720 to pass	\$175 list fee	12 months valid
---	-----------------------------------	--------------------------------------	--------------------------------------	---	--

WHERE THE MARKS ARE

Integration

Solution design

Evaluation and testing

Governance and risk

Stakeholders, lifecycle

Models and prompting

Developer enablement

19%

17%

16%

14%

14%

13%

7%

IF YOU REMEMBER NOTHING ELSE

- 1 Climb the pattern ladder only when the rung below cannot hold.
- 2 Least privilege means removal, not monitoring.
- 3 Retrieval degradation after a data change is an index problem.
- 4 Subjective quality is measured with rubrics and labeled sets.
- 5 Prompt changes are production changes: regression, then A/B.
- 6 Compliance is enforced before data crosses the boundary.
- 7 Route human review by confidence and consequence.
- 8 Present segmented, honest numbers, never a flattering average.
- 9 Observability captures traces, not just outputs.

ON THE DAY

Answer everything, unanswered scores zero · flag and return rather than stalling · read how many responses each item wants · find the stated constraint, then eliminate

Full sheet, study guide, notes, practice questions, and a timed mock exam: amey-thakur.github.io/CLAUDE-CERTIFICATIONS

**Amey Thakur**
github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS

ANTHROPIC
Facts drawn from the official Anthropic exam guide

The full sheet, with the reasoning behind every rule and a timed mock exam for this certification, is at amey-thakur.github.io/CLAUDE-CERTIFICATIONS



PART 3

Prepare

What to study, in what order, and how to practice without touching material you are not allowed to touch.

IN THIS PART

The official curriculum	19
A working study plan	20
The deck, one card at a time	21
One question, worked through	22
Two prompts worth keeping	23
Preparing	24



The official curriculum

Every Anthropic course, free on the public Academy, and the exam each one serves.

 OFFICIAL CURRICULUM

Every course, and the exam it serves

All of these are free on the public Anthropic Academy. No partner account is needed to learn the material; only the proctored exams require partner membership.

Claude platform

Claude 101	Associate
Claude Platform 101	Associate
Claude Code 101	Developer · Architect
Claude Code in Action	Developer · Architect
Introduction to Claude Cowork	Associate

Developer and integration

Building with the Claude API	Developer · Architect
Introduction to Model Context Protocol	Developer · Architect
Model Context Protocol: Advanced Topics	Architect Professional
Introduction to agent skills	Developer · Architect
Introduction to subagents	Developer · Architect

Deployment platforms

Claude with Amazon Bedrock	Architect
Claude on Google Cloud	Architect

AI Fluency

AI Fluency: Framework & Foundations	Associate
AI Capabilities and Limitations	Associate
AI Fluency for Builders	Associate
Roles: educators, pK-12, students, nonprofits, business, creative	Context
Teaching AI Fluency	Instructors

Per-course notes on what each is worth, what to watch for, and the order worth taking them in are in the repository.

 **Amey Thakur**
github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS

ANTHROPIC
Course catalog published on Anthropic Academy

Notes on what each course is worth and the order worth taking them in are at amey-thakur.github.io/CLAUDE-CERTIFICATIONS



A working study plan

Three weeks alongside a job: scope, then depth, then close.

 **A WORKING PLAN**
Three weeks, alongside a full-time job
Compress or stretch it to fit. The shape matters more than the calendar: scope, then depth, then close.

WEEK ONE
Scope

- Read the exam guide end to end
- Turn the blueprint into a red, amber, green checklist
- Work the prep courses for your exam
- Mark items green only when you could explain them

WEEK TWO
Depth

- Attack red items first, weighted by domain percentage
- Build the thing the guide asks you to build
- Read the official documentation for weak areas
- Drill one weak domain a day

WEEK THREE
Close

- Reread the exam guide
- Work the official sample questions and rationales
- Sit a timed mock and study the domain breakdown
- Run the system test on the exam machine

THE TWO RULES THAT MATTER MOST
Study by domain weight, not by what interests you · build something real, because the exams test judgment and judgment comes from hitting the tradeoffs yourself

Book the exam once the checklist is mostly green. Rescheduling is free until 24 hours out, so an early booking costs nothing and sets a deadline.

 **Amey Thakur**
github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS


ANTHROPIC
Method from the official preparation guidance

A progress tracker keeps this checklist for you, weighted by exam share, at amey-thakur.github.io/CLAUDE-CERTIFICATIONS



The deck, one card at a time

Every fact, domain weight, rule, and glossary term in this companion is also a flashcard. One hundred and ten of them, generated from the same source as everything you have read, so they cannot drift out of date.



Turn them in the browser, filtered by exam or by topic, at amey-thakur.github.io/CLAUDE-CERTIFICATIONS/guide/flashcards, or download the deck as a single tab-separated file that Anki, Quizlet, and RemNote all import directly. Use it for the facts that have to be automatic on the day: weights, codes, retake windows, and the terms the questions assume you know.



WHAT THE QUESTIONS ACTUALLY LOOK LIKE

One question, worked through

Every practice question in this repository is written the way the real items are: a scenario, one best answer, and three distractors that each fail a stated constraint. This one is from Architect Foundations. Read the scenario, commit to an answer, then read the reasoning.

Agentic architecture. Your support agent resolves most tickets but sometimes loops indefinitely, re-calling `lookup_order` on the same order. What is the correct structural fix?

- A. Raise `max_tokens` so the loop has room to finish
- B. Remove the `lookup_order` tool
- C. Add explicit loop-termination conditions: cap tool iterations, detect repeated identical calls, and route to escalation
- D. Lower the temperature

THE ANSWER IS C

Agent loops need engineered termination: an iteration budget, detection of repeated identical calls, and an escalation path for when the agent cannot make progress. The fix is structural, not a setting.

WHY THE OTHERS ARE WRONG

Raising the token budget (A) delays the symptom without bounding the loop. Removing the tool (B) breaks the resolution the agent exists to perform. Lowering the temperature (D) changes how the model words things and does nothing to iteration.

WHAT THE QUESTION IS REALLY TESTING

Not whether you know what an agent is. Whether you reach for a control when you see unbounded autonomy. Every distractor here is something a working engineer might genuinely try first, which is exactly how the real exam builds its options: plausible, and wrong for a reason you can name.

There are 319 more. Three hundred and twenty original questions, eighty per certification, split between topic practice and three timed mock exams each. Every one carries this kind of reasoning, not just a letter.

amey-thakur.github.io/CLAUDE-CERTIFICATIONS



Two prompts worth keeping

Claude is the most patient tutor you will get, but only if you ask properly. These two do more than any others. Type them as written, replacing the *coral italics* with your own details.

FIND OUT WHERE YOU ACTUALLY STAND

Run this before you study anything. It produces a ranked list, which beats working through the syllabus in order.

You are helping me prepare for the *Claude Certified Architect, Foundations* exam. Here is the published blueprint with the weight of each domain: *paste the domain table*.

Ask me one question at a time, twelve in total, spread across the domains in proportion to their weight. Scenario-based, one best answer, four options. Do not tell me the answer until I have committed to one.

After all twelve, give me a score per domain, the two domains where I am weakest in weight order, what specifically I misunderstood rather than which questions I missed, and what to study first.

UNDERSTAND A QUESTION YOU GOT WRONG

The highest-value minute in your preparation. A wrong answer you understand is worth more than ten right ones you guessed.

I got this practice question wrong. Question: *paste it with all four options*. I chose: *your answer*. The correct answer is: *the correct answer*.

Do not just restate the rationale. Tell me what rule the question is actually testing, why the option I chose is attractive and what constraint it fails, what would have to change in the scenario for my answer to become correct, and two other situations where the same rule decides it.

Seven more, including a study planner, an Architect scenario rehearsal, and one that reads a real score report and tells you the smallest gap to close, are at amey-thakur.github.io/CLAUDE-CERTIFICATIONS/guide/prompts



Preparing

A METHOD THAT WORKS

1. **Read the exam guide twice.** Once before studying, once after. It states plainly that the blueprint is the authoritative scope: anything outside it is not on the exam.
2. **Turn the blueprint into a checklist.** Mark each objective red, amber, or green, and study red first, weighted by the domain percentages rather than by interest.
3. **Build something real.** Every official guide asks for this. The exams test judgment in scenarios, and judgment comes from having hit the tradeoffs yourself.
4. **Use the official sample questions as calibration.** The rationales matter more than the answers: they show what the exam considers wrong, and why.
5. **Book once the checklist is mostly green.** Rescheduling is free until 24 hours out, so an early booking costs nothing and sets a deadline.

PRACTICING WITHOUT BREAKING THE AGREEMENT

Every candidate accepts a non-disclosure agreement covering exam questions, answer options, and scenarios, and it extends explicitly to online forums. Real exam content must never be shared or sought. There is also no official practice exam: the previous platform's was retired in the move to Pearson.

What works instead, and stays clean: generate your own questions from the published blueprint. The repository ships a bank of original questions with a shuffled, timed engine that runs in the browser or a terminal, and a set of Claude Code commands that quiz, drill, and plan against the blueprints.

The honest framing. Practice questions calibrate your knowledge against the blueprint. They are not exam items, and treating any question set as "the real ones" is both a policy violation and a poor strategy, because the item bank is confidential and rotates.



PART 4

Sit it

Booking, proctoring, and the day itself, then what the result means and what happens if it is not the one you wanted.

IN THIS PART

Exam day	26
Scoring and second chances	27
Policies and scoring	28



Exam day

Most failed sittings are logistics, not knowledge.

 **EXAM DAY**

The checklist worth keeping

Most failed sittings are logistics, not knowledge. Work through this the week before, then again on the morning.

A week before

- ☐ Run the OnVUE system test on the exact machine and network
- ☐ Send the required Pearson domains to your IT team
- ☐ Confirm your registration name matches your photo ID exactly
- ☐ Request accommodations if you need them, before scheduling
- ☐ Check your partner discount appears at checkout

The night before

- ☐ Reschedule free of charge only until 24 hours out
- ☐ Clear the desk, unplug the second monitor, book the room
- ☐ Charge the machine and locate your government-issued ID
- ☐ Reread the exam guide blueprint one last time
- ☐ Sleep. The paper rewards judgment, not cramming

On the morning

- ☐ Close browsers, Zoom, Teams, Outlook, Discord, remote access tools
- ☐ Close the Claude desktop app; OnVUE will not launch beside it
- ☐ Rerun the system test, even if it passed last week
- ☐ Start check-in early; it takes longer than you expect
- ☐ Eat something. Seat time is about 135 minutes

During the exam

- ☐ Answer everything: an unanswered item scores zero
- ☐ Flag and move on rather than stalling on one question
- ☐ Read how many responses each multiple-response item wants
- ☐ Find the stated constraint, then eliminate against it
- ☐ Report a faulty question afterwards; it never counts against you

Full registration and proctoring detail, including the complete domain allowlist and application shutdown list, is in the repository.

 **Amey Thakur**
github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS

ANTHROPIC
Requirements published by Anthropic and Pearson VUE

The complete network allowlist and application shutdown list are at amey-thakur.github.io/CLAUDE-CERTIFICATIONS



Scoring and second chances

How the score is built, what a failed attempt costs, and how the credential stays alive.

**SCORING AND SECOND CHANCES**

What happens after you click submit

How the score is built, what a failure actually costs, and how the credential stays alive.

HOW THE SCORE IS BUILT

100 – 1000 scaled range	720 to pass	criterion referenced, not a curve	per domain percent correct reported
-----------------------------------	-----------------------	---	---

Scaling equates forms of different difficulty. The domain breakdown does not decide the result, but it is the best study signal you will get.

IF YOU DO NOT PASS

1st attempt
wait 14 days

2nd attempt
wait 30 days

3rd attempt
wait 90 days

4th attempt
cap for 12 months

Every attempt costs the full fee, and your partner discount applies to retakes exactly as to a first sitting. Limits are per exam, so one failure never blocks a different track.

KEEPING IT ALIVE

Renew on time
A free, open-book, non-proctored assessment on Partner Academy, retakable as often as you need. Extends the credential another 12 months.

Let it lapse
The full proctored exam again, at the full fee. There is no partial credit for a credential that expired.

A disputed question never changes a pass or fail on its own; a confirmed faulty item that affected a result is remedied with a free retake.

**Amey Thakur**
github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS

ANTHROPIC
Facts drawn from the official Anthropic exam guides and policies



Policies and scoring

HOW THE SCORE IS BUILT

Each exam is criterion-referenced: you are measured against a fixed standard set by subject matter experts, not against other candidates. Results are reported on a scaled range of 100 to 1,000, with 720 required to pass, and scaling equates forms of slightly different difficulty.

Your score report also shows percent-correct per domain. Those percentages do not decide the result, but they are the most reliable study signal you will get, and they should drive the next plan.

RETAKES

Waiting periods grow with each failed attempt: 14 days after the first, 30 after the second, 90 after the third, with at most four attempts per exam in a rolling 12 months. Each attempt costs the full fee, and partner discounts apply to retakes as to first sittings.

VALIDITY AND RENEWAL

A credential lasts 12 months from the day it is earned. Renewing before it expires is free: an open-book, non-proctored assessment on Anthropic Partner Academy, retakable as often as needed, extending the credential another 12 months. The credential itself arrives as a digital badge from Credly, by email, shortly after a pass. Let it lapse and the full exam must be passed again at full fee.

APPEALS AND FAULTY QUESTIONS

Decisions can be appealed to Pearson within 14 days. Separately, anyone may report a question that looks wrong, ambiguous, or out of scope; reporting never counts against you, and a confirmed faulty item that affected a result is remedied with a free retake rather than a changed score.

Eligibility. Certification is currently open to people at Claude Partner Network organizations, registering with a recognized company email. Domain record changes take 7 to 10 days, so resolve any email problem well before you plan to sit.



PART 5

Keep going

You now have everything the exam asks of you. What follows is where to find the parts that keep moving.

IN THIS PART

The words the exams assume you know	30
The 22 certificates behind this companion	31
The repository behind it	32
Where to go next	33



REFERENCE

The words the exams assume you know

The exam guides use these terms without defining them, and questions are written on the assumption that you read them the way the program does. 24 terms.

Blueprint. The exam content outline in each exam guide: the domains measured, their weights, and the objectives items are written against. The authoritative scope of the exam

Domain weight. The approximate share of scored items drawn from a domain, for example 33.1% for the Developer exam's Applications and integration domain

Task statement. A numbered capability within a domain, used in the Architect Foundations guide, that items are written against

MQC. Minimally qualified candidate: the profile of the weakest examinee who should still pass, which the passing standard is calibrated to

Criterion-referenced. Scored against a fixed standard rather than against other candidates. There is no curve

Scaled score. The reported score on a 100 to 1,000 range, equated so results are comparable across exam forms of slightly different difficulty

Cut score. The passing threshold: 720 scaled, for all four exams

Item. An exam question. Multiple-choice items have one best answer; multiple-response items state how many responses to select

Scenario-based exam. The Architect Foundations format: 4 scenarios drawn from a published bank of 6, each framing a set of items

Exam form. A specific assembled version of an exam; forms vary slightly, which scaled scoring corrects for

NDA. The non-disclosure agreement accepted before the exam starts: exam content may not be shared, reproduced, or discussed, including in study groups and forums

Pearson VUE. The company that delivers the exams, online and at test centers, since June 30, 2026

OnVUE. Pearson's online proctoring application, run on your own computer under remote supervision

Test center. A physical Pearson location with provided equipment, the alternative to online proctoring

Credly. The digital badging platform that issues the shareable credential after a pass

CPN. The Claude Partner Network, whose member organizations are eligible for certification

Partner tier. An organization's CPN level: Registered pays list price; Select, Preferred, and Global Premier receive 50% off exams

Prep path. The curated course sequence on Partner Academy for a specific certification

Exam guide. The official PDF for each certification: blueprint, sample questions, policies, and scoring. The single most important study document

Renewal assessment. The free, open-book, non-proctored assessment on Partner Academy that extends a credential 12 months when passed before expiry

Lapse. Letting a credential expire without renewing; regaining it requires the full exam at full fee

Appeal. A formal challenge to a certification decision or result, filed with Pearson within 14 days

Question report. Flagging a possibly faulty item to Pearson; separate from an appeal and never counted against you

Accommodations. Testing adjustments for documented needs, requested through Pearson and approved before scheduling



RECEIPTS, NOT CLAIMS

The 22 certificates behind this companion

Every course in the official Anthropic Academy curriculum, completed before this companion was written. Each certificate carries its own Skilljar verification record, and 19 of them were also issued as a digital completion badge on Claude Academy, verifiable on Anthropic's own domain.

THE 19 CLAUDE ACADEMY BADGES



ALL 22 COURSES

CLAUDE PLATFORM

- Claude 101
- Claude Platform 101
- Claude Code 101
- Claude Code in Action
- Introduction to Claude Cowork

DEVELOPER AND INTEGRATION

- Building with the Claude API
- Introduction to Model Context Protocol
- Model Context Protocol: Advanced Topics
- Introduction to agent skills
- Introduction to subagents

DEPLOYMENT PLATFORMS

- Claude with Amazon Bedrock
- Claude on Google Cloud

AI FLUENCY

- AI Fluency: Framework & Foundations
- AI Capabilities and Limitations
- AI Fluency for Builders
- AI Fluency for educators
- AI Fluency for pK-12 Educators
- AI Fluency for students
- AI Fluency for nonprofits
- AI Fluency for Small Businesses
- AI Fluency for Creative Work
- Teaching AI Fluency

Every certificate, its PDF, and its verification link is at amey-thakur.github.io/CLAUDE-CERTIFICATIONS/certificates. The courses are free to anyone on the public Academy.



BEYOND THIS COMPANION

The repository behind it

This companion is a snapshot, printed on a date. The repository is the living version: maintained, link-checked every week, and updated when Anthropic changes the program. These are the things paper cannot do.

PRACTICE, NOT JUST READ

A bank of original practice questions with a shuffled, timed engine that runs in your browser or your terminal. It reorders questions and options on every run, so nothing can be memorized by position, and it scores you per domain the way the real report does. Three timed mock exams per certification, each with questions that appear nowhere else, so a second sitting still measures something.

STUDY THE WAY YOU ALREADY STUDY

A flashcard deck covering every fact, domain weight, rule, and term, in a format Anki, Quizlet, and RemNote all import directly. A progress tracker that weights each domain by its share of the exam, so the number tells you how much of the paper you can actually answer. Anthropic's own videos, including the complete AI Fluency course, arranged by the exam each one serves and playable in the page.

A COACH THAT KNOWS THE BLUEPRINT

Open the repository in Claude Code and six commands are waiting: run a diagnostic to find weak domains, drill one of them, sit a timed mock, build a study plan or a week of sessions, and interpret a real score report into a decision.

THE OFFICIAL DOCUMENTS, KEPT CURRENT

Every official exam guide and policy PDF, mirrored with its source URL and the date it was last checked, so you can verify anything here against the original in one click. A weekly job re-checks every link, and a script re-downloads the documents and reports what changed.

THE PARTS THAT GO STALE FASTEST

Fees, discounts, question counts, and the migration-era rules change. The repository carries dated facts and a maintenance routine designed to catch them. When something here disagrees with the official exam guide, the guide is right, and the repository is where the correction lands first.

OTHER CANDIDATES

A discussion space for preparation questions and exam experiences, under one firm rule: real exam content is never posted. If this companion helped, the most useful thing you can do is answer someone else's question there.

Everything is free and open source. No account, no paywall, no email capture. Take it, use it, fork it, and send it to whoever is studying next.

github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS



Where to go next

FOUR PLACES TO GO

github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS

The repository: practice engine, flashcards, progress tracker, and the mirrored official documents.

amey-thakur.github.io/CLAUDE-CERTIFICATIONS

The same material as a searchable website, with the practice engine in the browser.

anthropic.skilljar.com

The official courses, free, without a partner account. Registration and the exams run through Anthropic Partner Academy and Pearson VUE.

github.com/Amey-Thakur/CLAUDE-CERTIFICATIONS/discussions

Questions about preparing, answered in the open, so the next candidate finds the answer instead of asking again.

IF THIS HELPED

Share it with whoever is studying next. Corrections are welcome as issues, and questions about preparation belong in the repository's discussions.

If it saved you time, starring the repository is the whole marketing budget: it is how the next candidate searching for this finds it instead of a braindump site.

Whatever you are preparing for, the path is the same: read the blueprint, close the gaps honestly, build something real, and book the date. You do not need permission or a better moment.

– Amey Thakur

Facts drawn from the official Anthropic exam guides and program pages, verified against them on publication. A community study resource, not affiliated with or endorsed by Anthropic. Claude and Anthropic are trademarks of Anthropic PBC.

